

AI Assisted Work and Accepted State

Theory, evidence, limitations, and the next phase

Project by Jake Tiller | Evidence review 5 September 2026

The project in plain language

This project asks how AI-assisted work can become a dependable operational record without letting a plausible answer authorize itself, hide an uncertain effect, or erase a previously accepted result.

The common thread is simple: let models propose and attempt bounded work; require explicit authority and qualified evidence before advancing accepted state; retain failures so the next attempt can address a specific defect.

Current assessment: this is an applied systems-engineering research program with working reference implementations, finite tests, and useful negative results. It is more developed than an idea alone. General cyber-defense effectiveness, physical stability, and economic superiority remain unestablished. [L1 pp.2,8,11; L2 pp.1-3; G1]

What can be said confidently

- Clutch's public release checks passed again during this review at its exact commit. These are finite software checks, including invalid transitions and recovery routes.
- Existing local evidence was recovered for later Clutch studies, BespokeNode, the State Runner, and dREIT's HTTP and economic experiments. Reading those records is distinct from rerunning their original experiments.
- The strongest recurring contribution is explicit separation: proposal, permission, execution, effect, observation, acceptance, and later outcome have different meanings.
- Negative results matter: strong input reduction coexisted with failed delivery, and additional structure sometimes underperformed simpler controls.

Where to go next

Use one home page and one active task. First make one failed transition and its recovery independently replayable at zero API cost. Then freeze a small experiment that distinguishes replication of fixed code from improvement across repaired versions. Additional credits buy observations and development opportunities; they do not guarantee a favorable plateau.

The shareable claim is **a testable architecture for preserving governed state around uncertain AI work**, with security and efficiency as application hypotheses. No claim here certifies universal safety, originality, or a positive return on spending.

One project with several functions

The names describe parts of a common workflow. Their results should remain separate so a pass in one component cannot silently validate another. [L1 pp.4-7,20; L3 pp.1-5]

Component	Plain job	Evidence boundary
Open Mirror	Explore a question through analogies, find where they break, return to literal language.	Written exploration workflow; current version labels conflict. No truth or permission gate.
Pin Check	Check whether a label or number follows a stated source and rule.	Reported bounded method demonstration; independent replication remains open.
Delta Atlas	Compare terms by function and expose mismatched definitions or views.	Public data and inspection tools; no universal semantic validator.
Calibration Ledger and STP	Record proposals, evidence, decisions, effects, and corrections with explicit transition rules.	Narrow profiles implemented; full lifecycle and broader guarantees proposed.
Clutch	Transfer approved task context to a worker and define bounded recovery routes.	Executable protocol and evaluation kit; delivery and savings are separate empirical questions.
BespokeNode and State Runner	Compile local constraints; preserve and rebuild accepted operational records.	Local candidate and single-writer tests; real host enforcement and power-loss durability unqualified.
dREIT	Apply these boundaries to service obligations, spending exposure, dispatch, and reconciliation.	Historical local HTTP/economic lab; financial v3 remains proposed.

The working sequence

Question -> explicit hypothesis -> bounded task -> attempt -> evidence -> accept or hold -> recover or revise.

Open Mirror belongs before a claim is promoted. Clutch belongs at a handoff. The ledger preserves the history. State Runner reconstructs the operational view. dREIT is an application of the architecture, not a separate proof of the architecture.

The two supplied Clutch graphics are useful explanations. The earlier three-state graphic emphasizes a single approved attempt. The V2 graphic adds finite repair routes and last-accepted-state preservation. They represent different levels of the design; neither is evidence that a host enforces the pictured boundaries. [L6-L7; G1 host-integration]

The theory and its realistic meaning

A small operational core

Let A be the last accepted record, C a candidate, and G the applicable admission rule. The intended transition is:

A_{next} = C if G passes and the authorized acceptance is recorded; otherwise A_{next} = A.

G must refer to exact content, current policy and lineage, required evidence, a valid event history, the declared effect/recovery conditions, and the relevant acceptance decision. A simple Boolean formula is only meaningful when those inputs have defined sources and cannot be forged or bypassed. [L1 p.8; G1]

A second useful idea concerns compact context. If two histories become identical after compression but require different decisions, the compressed representation is insufficient for that decision. This can be tested in a finite domain. It does not establish that all task-relevant meaning has survived. [L1 p.9; L2 pp.17,21,31-33]

Is this a dynamic state system

Yes, in the ordinary software sense: events change records, admissible actions depend on current state, and feedback can change the next candidate. It is an architecture and research hypothesis about governed transitions.

It is not yet a demonstrated physical dynamical theory. Physical stability requires a defined state space, inputs, observations, disturbances, timing, and a suitable stability criterion. A neutral software label is not necessarily a safe physical condition. A dispatched action may already have affected the world even when the local record holds. [L2 pp.4-8; L3 p.9; P3-P4]

Cyber-defense relevance

Exact authority, restricted interfaces, stale-state rejection, replay checks, explicit uncertainty, and bounded recovery are relevant control mechanisms. NIST's cyber-resilience framing concerns the ability to anticipate, withstand, recover, and adapt. This project addresses parts of that problem; incident reduction and compromised-host survival require separate evidence. [P2]

A gate over caller-supplied JSON checks a representation. A production defense also needs authenticated actors, enforced permissions, trustworthy observations, durable state, and protection against direct access around the gate. Agreement between two reducers cannot establish these conditions. [G1; L2 pp.25-26]

What the evidence currently supports

REPRODUCED means an existing offline check was rerun in this review. SOURCE CHECKED means an exact local or public record was inspected and its arithmetic or identity checked; the original experiment was not rerun. These labels describe this review, rather than inheriting the source author's TESTED label.

Evidence	Result	Supported reading
Clutch public v0.4.0-rc.1	REPRODUCED: release suite passed; 30 conformance cases; 48 Fuse cases, including 29 hostile cases.	Finite local conformance and release-byte checks. No live routing or security guarantee.
Public STP packet	REPRODUCED: 47 synthetic cases matched expected results; Python/JavaScript projections agreed.	Finite replay of concurrency, effects, capacity, instruments, recovery, and economy. Full release suite not rerun.
Clutch six-round pilot	SOURCE CHECKED: 6 accepted vs 5 direct; 56.5% fewer raw tokens; 15.1% longer elapsed time.	Descriptive signal. Registered comparison gate allowed no efficiency winner.
Fixed v2 study	SOURCE CHECKED: 15 tasks; about 95.9% input reduction. Stable completion 9/15 cards vs 15/15 history; switched 13/15 vs 10/15.	Average-score noninferiority did not preserve strict completion.
Factorial coherence	SOURCE CHECKED: 192 cells; collapse 12/96 cards vs 35/96 history.	Exploratory structure/context signal; primary switch question inconclusive.
Repeatability	SOURCE CHECKED: 2/42 deliveries in each route; \$43.4289 total accounted spend.	Completed-design delivery failure. No favorable reliability or savings conclusion.
BespokeNode Candidate 004	SOURCE CHECKED: 256 Boolean assignments; 103,788 stress cases per reported pass; one Windows reproduction receipt.	Bounded local behavior. Reused dependencies limit reproduction independence.
State Runner	SOURCE CHECKED: 12,000 valid events plus 12 duplicate attempts; no recorded silent loss or accepted-state drift.	One local single-writer run. No power-loss, concurrency, or production guarantee.
dREIT historical labs	SOURCE CHECKED: 41 HTTP scenarios and three negative controls across two passes; economic workload records recovered.	Local model/HTTP evidence. v3 financial controls are not implemented by these results.

Sources: G1-G2; L1 pp.11-18; R1-R6. Counts are not additive across suites, seeds, phases, or revisions. Repeated transitions are not necessarily independent samples. The report does not turn large test counts into a percentage of real-world coverage.

Failures that changed the design

These are research assets because they identify a claim to narrow or a check to strengthen. Retaining a failure is useful when it changes a decision; failure alone does not prove that improvement will follow.

Observed or recorded problem	What it teaches	Response or next step
Compact cards met an average score margin but stable delivery fell from 15/15 to 9/15.	Average quality can conceal an unacceptable completion loss.	Gate strict delivery separately from average score.
Both repeatability routes delivered 2/42.	More context and governed cards both failed in this domain.	Preserve the frozen result; diagnose failure families before another outcome claim.
A 376-byte projection passed an incomplete test domain and then failed corrective cases.	An optimizer can exploit holes in the evaluator.	V6 retained task/authorization distinctions; 431 bytes was sufficient only in the registered domain.
Two fields in one deployment receipt disagreed about evidence and authority while schema checks passed.	Valid shapes do not imply cross-field consistency.	PROPOSED: relationship checks and an explicit vocabulary migration; defect remains open.
RELATED absorbed 7/8 Open Mirror slots; harder rules shifted traffic toward CONTESTED.	Wording pressure can move label gaming elsewhere.	Test decision changes and negative examples; retain UNDEFINED/NO INCREMENT outcomes.
P1-P3 refusal scaffolds had unsupported claims; a larger-model control recorded 0/151.	Structure is not automatically better; model and scaffold effects were confounded.	Keep controls and the missing raw-data limitation.
Three software telemetry channels remained coupled; 6/943 work samples looked idle under the reported thresholds.	Multiple channels do not automatically constitute independent witnesses.	State sensor and adversary assumptions; do not call a channel unspoofable or a work oracle.
dREIT local HOLD could coexist with a partial external effect.	Preserving accepted records does not reverse the world.	Reconcile the destination effect; name supported compensation or finality conditions.

Sources: L1 pp.9,12-15; L2 pp.22-25; L4 changelog; G1 projection evidence; G2 power Phase3; G3; L3 p.9; R5.

Resilience means preserving a recoverable state and a useful next action. Use a bounded repair, rerun the relevant check, and stop when the evidence is missing or the budget is exhausted. A stop can be correct even when no improvement is earned.

The plateau hypothesis

On 5 September 2026, Jake reported that limited API funding may have shaped inadequate test constraints, and proposed that another \$200 might have allowed a later plateau. He clarified that delivery, cost per success, and failure/instability could stabilize together or in phases. His point concerns the adequacy of the experimental resource envelope, not merely interrupted calls. This remains an owner-stated hypothesis. [U1]

Three measurements instead of one metaphor

Dimension	Measurement to freeze	What would count against the useful hypothesis
Delivery	Strict accepted deliveries divided by all scheduled, eligible tasks; report missing runs separately.	A stable rate below the required quality floor, or no improvement on held-out tasks.
Economics	Total route cost, including failed attempts and review, divided by accepted deliveries.	Stable cost above the comparator; zero deliveries; hidden labor or missing telemetry.
Failure and instability	Named failure counts, critical violations, recovery success, and between-window variability.	Persistent critical failures, rising variability, or recovery costs that exceed the benefit.

Distinguish failure types and volatility: a binary failure rate would only duplicate delivery's complement.

Two different curves

Fixed system: more trials estimate unchanged behavior. Testing itself does not repair it.

Repair trajectory: record interventions and versions, then evaluate on untouched tasks. Do not pool changing versions into one replication rate.

Money as a feasibility gate

If necessary resources exceed the budget, success through that plan may be unreachable. Adequate funding does not guarantee success. The threshold depends on the task, implementation, output/context limits, observation, review, and recovery; it is not a universal dollar threshold or physical law.

A completed test can still poorly represent its intended operating regime. The factorial study's 24 responses at a 6,000-token ceiling justify examining resource sensitivity. They do not establish why that limit was chosen or whether raising it would fix delivery. [R3]

What the stopping record does establish

The repeatability run completed 301 scheduled calls, spending \$43.4289 inside its \$64 ceiling. Its 2/42 result stands under the frozen constraints. Completion does not establish that the constraints were adequate, what the account balance was, or what a larger resource envelope would achieve. [R1]

Before paid continuation, vary a named resource constraint in matched tasks while holding other conditions fixed. Separate this resource-response experiment from a repaired-version trajectory. Define plateau windows, tolerances, minimum useful quality, and a spending stop in advance. More resources could reveal improvement, an unacceptable plateau, or continuing instability; current evidence does not select among them.

Economic efficiency and physical limits

The economic hypothesis is conditional: preparation and review may pay for themselves if bounded execution saves enough work while delivering an acceptable artifact. Cheaper calls, fewer prompt bytes, faster elapsed time, and lower total operating cost are different measurements. [L1 p.10; G1 evaluation; P5]

Total route cost = preparation + all attempts + verification + tools + measured human labor + measured recovery.

Cost per accepted delivery = total route cost / number of accepted deliveries.

If there are no accepted deliveries, the cost-per-delivery ratio is undefined. If a required cost component is missing, report an incomplete subtotal and UNKNOWN total. Do not divide only the cost of successful attempts by successes and present that as workload efficiency. The repeatability record's successful-route costs are expressly diagnostic, not that full-workload measure. [R1]

An unfavorable result worth keeping

In dREIT's synthetic 20-job burst workload, contract-aware routing cost 158 fictional cents versus 156 for static fallback. Mean wait per attempt was 26 versus about 13.36 logical ticks; provider contribution was 52 versus 62 fictional cents. That is about 1.3% higher workload cost, 94.6% longer mean wait, and 16.1% lower provider contribution. These are synthetic model units, not invoices or wall-clock latency. [L3 p.8; R6]

The implemented controls can have value even when this comparison loses, but avoided loss and operational benefit must be measured. They cannot be invented to make the arithmetic favorable.

What physical metaphors can and cannot carry

A clutch suggests disengagement; a reservoir suggests backlog; a plateau suggests diminishing change. Each can generate a testable question. The software inherits none of the physical mechanism simply because the shape resembles it.

For observation, settlement, and recovery queues, keep units and arrivals separate for each lane. The earlier STP candidate uses a common arrival/unit formulation; the later working paper corrects this to lane-specific quantities. Preserve that as a documented revision, not a silent alias. [L5 section 8; L2 Appendix A.6, pp.30-31]

Runtime security requires actual mediation before an effect. Crash recovery requires tested persistence and destination semantics. Energy efficiency requires energy measurement with a declared boundary; fewer tokens alone do not establish fewer joules. The project presently offers useful hypotheses and local checks at these interfaces.

Contribution and intellectual lineage

The reviewed materials support a concrete account of Jake's work: selecting the problem, defining authority and acceptance boundaries, directing experiments, retaining counterexamples, and shaping the implementations and explanations. They explicitly disclose substantial AI assistance with research, drafting, coding, testing, and review. [L1 p.23; L2 pp.27-28; L4 Credits]

The distinctive synthesis visible here

- **No silent promotion across boundaries.** An analogy does not become evidence; an answer does not become authority; a receipt does not become external truth.
- **Decision-preserving context.** Compression is assessed by whether histories requiring different actions become indistinguishable, rather than by size alone.
- **Recovery with retained costs and failures.** Failed work can supply evidence without replacing accepted state or disappearing from the economic account.
- **A falsifiable repair history.** Delivery regressions, evaluator shortcuts, identity collisions, and contradictory receipts are used to narrow the design.

These are specific contributions of emphasis, integration, implementation, and operational method visible in this corpus. They are not findings that no one else has combined similar ideas.

Established neighbor	What should be credited	Project-specific application
Event sourcing	Event histories, projections, and compensating records.	Typed AI proposal/authority/observation/acceptance records.
Least privilege and runtime assurance	Restricted authority; trusted mediation and fallback.	Approved task cards and explicit non-promotion conditions.
Feedback control	State, sensing, disturbance, feedback, and stability analysis.	Proposed governance and recovery measurements.
Cost-aware model cascades	Trading model configuration against cost and performance.	Human-governed context handoff and delivery-conditioned accounting.
Analogical structure-mapping	Comparing relations across domains.	Open Mirror's baseline, break-point, and decision-change workflow.

Primary sources P1-P6 support these neighboring mechanisms; none endorses this project.

What can establish a credible authorship record

Preserve dated drafts, exact commits, source citations, AI contribution notes, and the sequence from hypothesis to counterexample to revision. A digest proves identity relative to an anchor, not authorship or priority. This targeted review cannot certify absence of plagiarism or worldwide novelty. The existing attribution and repair trail is a stronger account of the work than an unsupported claim of being unprecedented.

Strengths and operating constraints

The strengths below are visible in the artifacts. They are assessments of the work, rather than personality scores.

Strength	Evidence	How to use it
Translating ideas into constraints	Exact acceptance, authority, and projection conditions.	Require every new metaphor to yield one literal question or test.
Willingness to keep disconfirming results	2/42 repeatability, negative economics, scaffold controls.	Put the strongest counterexample beside the strongest claim.
Finding hidden distinctions	Dispatch vs effect; cap vs duplicate obligation; replay vs truth.	Build small fixtures that distinguish two easily confused states.
Concern for recovery and human control	Holds, finite repair routes, fix logs, accepted baseline.	Make interruption and return part of normal operation.

The main bottleneck is breadth: the work spans semantics, distributed state, model evaluation, host security, economics, and service settlement. Each requires different evidence. Many new names and versions increase the effort required to resume or explain the project.

Other concrete weaknesses are uneven raw-evidence availability in public packages, version ambiguity, occasional absolute language, and reliance on common specifications or operators when discussing independence. The system also risks becoming too expensive to maintain if every note requires a full formal packet.

A practical rhythm for limited bandwidth

Keep one active task and a parking list. A normal session has one question, one artifact, one check, and a stopping condition. Match the record to the consequence: a speculative note needs a source and a status; a consequential action needs the fuller contract.

At the end of a session, leave four lines:

Current version and last accepted result:

What changed and where the evidence is:

Next physical action:

Stop or ask for help if:

On a low-capacity day, read those four lines and complete one small evidence check. On a better day, finish the active card before opening another workstream. Measure resilience by how reliably work can resume and reach a checked result, rather than hours of uninterrupted effort.

A home for the files

The accompanying Project_Home folder supplies an index, preserved copies of the seven supplied artifacts, claim and failure registers, an active task, and proposed GitHub edits. Existing research folders remain in place and are linked through a private source register.

Location	Put here	Rule
00_HOME.md	Purpose, current claim ceiling, next task, navigation.	First file opened after an interruption.
01_SOURCES	Original attachments and source identities/locators.	Preserve original bytes; record corrections separately.
02_RESEARCH	Editable report, hypotheses, terminology, and prior-art notes.	Label owner hypotheses separately from findings.
03_EVIDENCE	Run register, failures, validation receipts, detailed reviews.	Link a claim to a source, version, result, and limit.
04_NOW	One active card, restart note, parking list.	One task in progress; no automatic retry loop.
05_GITHUB	Concrete proposed README and cleanup changes.	Reviewable proposals; no remote changes performed.

A small record instead of a new bureaucracy

Every claim needs an ID, a sentence, evidence status, source locator, limit, and next falsifier. Every failure needs expected vs observed behavior, the affected version, consequence, and a repair or deliberate hold. A source register needs full digest and location. These three records carry most of the organizational load.

Use one current pointer, not several files named final. A newer date does not automatically supersede a different protocol. Accepted records retain their identity; a correction names the earlier record and explains what changed. Move a current pointer only after the intended acceptance decision.

Version issues already identified

Open Mirror's filename/header says v1.1 while its changelog reaches v1.1.2 UNTESTED. STP v1.2 remains a candidate and references a Ledger identity gap; Ledger v9 does not automatically resolve it. The later paper's queue formulation differs from the candidate. Clutch graphics, public release v0.4, and later local evaluations need a version map rather than a single blended status. [L1; L2; L4-L7]

The PDF is the single shareable narrative. The source archive and private locators are working material. Share a minimized reproduction package separately when a collaborator needs exact artifacts.

GitHub attention in priority order

Seven public repositories were found through the connected account. The public search succeeded even though the connector's installation/repository listing returned empty results. This review therefore does not establish a complete inventory of private repositories.

Repository	Current role	Recommended attention
Resilience-Ledger	Main public case study and STP evidence.	Use as the existing portfolio entry point; resolve the intended relationship of open paper PRs #14 and #15.
clutch-skill	Most cohesive public packaged protocol.	Keep its code and checks separate; link the later evidence chronology and precise empirical limits.
typed-refusal-harness	Retained negative/aggregate model study.	Label it an aggregate archive; identify missing original corpus/runs and exploratory statistics.
the-stable	Instrumentation and agent experiment lineage.	Correct README's statement that no real agent has played; findings/logs describe later real-agent trials.
the-crosswalk	Decision aids for coding, learning, directing AI, or hiring.	Identify this practical guidance role; distinguish it from metaphor crosswalks.
GroundingHarness	Earlier lineage, currently lacking README.	Restore a minimal archival description and pointer.
AI-Governance-Ledger	Earlier lineage, currently lacking README.	Restore a minimal archival description and pointer.

Three changes to do first

- 1. Establish one entry point.** Add the prepared portfolio map to Resilience-Ledger. Explain the project question, what runs today, strongest negative result, and one next contribution. Preserve existing repository history.
- 2. Reconcile review state.** Resilience-Ledger #14 is a draft v3 paper packet; #15 is a later v4 packet and has passing CI at the inspected head. Choose which packet is intended to supersede the other and compare exact artifacts before any merge. Passing CI is not paper acceptance. Both remained open at inspection. [G2]
- 3. Repair misleading status.** Add the two missing README pointers, update the-stable's stale statement, and narrow physical-witness language such as unspoofable or fully mapped. Sensor readings require a stated threat and measurement model. [G2; G4]

A focused code finding is also recorded: the lexicon engine rejects a negated control alias but can still credit an equivalently negated bridge cue. A proposed local patch and regression check accompany the audit; this matters because mentioning a control as absent must not count as its presence. [G2 audit]

No repository was renamed, archived, merged, or changed remotely during this review. Concrete drafts are in 05_GITHUB, so the next decision concerns visible text rather than an abstract cleanup request.

The next phase without new API spending

One active outcome

Make one accepted-state preservation example easy for another person to inspect: an allowed transition advances state; a stale or unauthorized candidate cannot advance it; the failed attempt remains available; a declared valid recovery reaches the intended result.

Start from the already-rerun public Clutch release or the recovered Candidate 004 evidence. Select one implementation for the packet. Avoid combining them into a new framework before that example is clear.

Three bounded sessions

Session 1: recover the exact baseline. Open the source register, select one commit and one failure/recovery fixture, and record the command and expected output. Done means the evidence can be located without reading old chats.

Session 2: make the example reviewable. Include a minimal input, the relevant rule, expected accepted-state identity before and after, observed output, and one counterexample. Explain the trusted host assumptions. Done means a reviewer can tell what passed and what remains outside the test.

Session 3: freeze the continuation design. Name the potentially inadequate resource constraint, three plateau metrics, task unit, baseline, matched resource levels, separate repair/version blocks, minimum useful quality, and cost stop. Done means the experiment can distinguish resource insufficiency from a defect, with an interpretable negative outcome.

When further spending becomes useful

First verify local checks and telemetry without model calls. Then estimate an affordable task count from actual per-task costs, with reserve for failed attempts and analysis. Treat the hypothetical \$200 as an outer resource limit, not a commitment or a guarantee of statistical power. A shorter instrumentation study may be justified before a larger efficacy study.

Freeze the sample and stopping rule before outcomes are examined. Do not repeatedly extend the sample only because a desirable plateau has not appeared. Record a budget stop as censored future observation, keep completed failures, and distinguish development progress from confirmatory evidence.

The most valuable outside contribution is one exact reproduction or counterexample under a different operator or host. More AI agreement from the same prompt and fixture corpus is helpful review, but it does not supply independent external validation.

Sources and review boundaries

Full source digests and private local locators are in Project_Home/01_SOURCES/source-register.csv. Short identifiers below are navigation aids, not replacements for full digests. Review date: 5 September 2026.

ID	Artifact and scope
L1	Calibration_Ledger_v9.pdf, 23 pages, dated 27 August 2026. SHA-256 a593ca4a2fa8... All pages reviewed.
L2	From-Model-Output-to-Accepted-State.pdf, 36 pages, owner-review paper dated 15 August 2026. SHA-256 483c4f320d58... All pages reviewed.
L3	dREIT-v3-Review-Package.zip. SHA-256 bc8a5eb986ca... Five files; four internal manifest entries verified. Ten-page PDF and protocol/JSON reviewed.
L4	Open-Mirror-v1.1.md. SHA-256 2f187067f372... Complete file reviewed; internal version ambiguity retained.
L5	State_Transition_Protocol_v1.2_Candidate.docx. SHA-256 8171dd3fb0e7... Complete OOXML body reviewed; section locators used, not unverified page numbers.
L6-L7	Clutch V2 graphic (extensionless PNG) and Clutch-Skill Graphic.jpg. Visually inspected as explanatory graphics.
G1	github.com/JakeTOpenSource/clutch-skill at e56bf1e59952acbabf5f3e61a8984d0e818e0791. Existing release checks rerun on Windows, Node 24.19.0, exit 0.
G2	github.com/JakeTOpenSource/Resilience-Ledger; main and paper PRs #14/#15 inspected. Detailed pinned paths and heads are in the GitHub audit.
G3-G7	typed-refusal-harness; the-stable; the-crosswalk; GroundingHarness; AI-Governance-Ledger, all under JakeTOpenSource. Public metadata, relevant artifacts and limited offline checks reviewed.
R1-R3	Recovered Clutch repeatability analysis/state/records; fixed v2 research summary; factorial handoff. Local source checks, not new provider runs.
R4	Recovered Candidate 004 finite-transition and State Runner records and Windows attestation for commit 4c03cd7a5b8badc7f15d38a5d1f33e7713b6e11c.
R5-R6	Recovered http-effect-lab-001 and digital-reit-adapter-001 reports/results. Identity and record inspection; no live effects or payments rerun.
U1	Jake's statements on 5 September 2026: funding may have shaped flawed test constraints; money can gate feasible success; three plateaus may occur together or in phases. Owner-reported constraint and hypothesis.

Local attachment instructions were analyzed as project material. They did not authorize new experiments, spending, publication, or changes to the host. Reviews by several AI agents are coordinated checks, not independent research-team replication.

Primary references and contribution statement

The comparison below is targeted rather than exhaustive. It establishes relevant prior mechanisms and limits; it does not assess patent novelty, certify originality, or claim endorsement.

P1. Microsoft Azure Architecture Center, Event Sourcing pattern. Event histories and reconstructed projections are established design techniques with concurrency, versioning, and operational tradeoffs.

<https://learn.microsoft.com/en-us/azure/architecture/patterns/event-sourcing>

P2. NIST SP 800-160 Volume 2 Revision 1, Developing Cyber-Resilient Systems, 2021; NIST least-privilege definition. Engineering resilience and restricted authority provide context for the proposed security applications.

<https://csrc.nist.gov/pubs/sp/800/160/v2/r1/final>

https://csrc.nist.gov/glossary/term/least_privilege

P3. Carnegie Mellon Software Engineering Institute, Simplex report, 1996. Trusted mediation and an adequate fallback are established runtime-assurance ideas. A local HOLD alone is not a physical fallback controller.

https://www.sei.cmu.edu/documents/1146/1996_005_001_16463.pdf

P4. Astrom and Murray, Feedback Systems author materials. State, observation, feedback, and stability require explicit mathematical and physical assumptions.

https://murray.cds.caltech.edu/CDS_101/110%2C_Fall_2015

P5. Chen, Zaharia and Zou, FrugalGPT, 2023. Cost-aware model combinations and cascades precede this project. Clutch's economic hypothesis must be evaluated against appropriate controls.

<https://arxiv.org/abs/2305.05176>

P6. Gentner, Structure-Mapping, 1983. Analogical comparison of relational structure is established prior work. Open Mirror contributes a particular workflow and testable grading rules.

<https://groups.psych.northwestern.edu/gentner/papers/Gentner83.2b.pdf>

A concise description for collaborators

Jake Tiller develops and tests workflows that keep AI-generated proposals separate from permission, execution, observation, and accepted records. The work combines explicit transitions, retained uncertainty, bounded delegation, and failure-driven repairs. It draws on established systems engineering and discloses substantial AI assistance. Its strongest results concern inspectable local behavior and counterexamples; broader security, delivery, economic, and plateau claims remain subjects for testing.

This synthesis and evidence organization were prepared with OpenAI Codex, including coordinated AI-assisted source reviews. Historical authorship and contribution statements are attributed to the project materials. The review created no new model-efficacy result and made no paid provider calls.